

Decomposing the GP Advantage in Offline MBO

Anonymous submission

Abstract

A common intuition in offline model-based optimization (MBO) is that Gaussian-process lower confidence bound (GP-LCB) outperforms neural-ensemble LCB at low dimension because the GP posterior is better calibrated. This comparison is confounded: the two pipelines differ in *both* the surrogate class and the acquisition optimizer. We run the first controlled surrogate $\times$ optimizer decomposition—{deep ensemble, exact GP, sparse variational GP} $\times$ {gradient ascent, perturbation, CMA-ES}—on 7 synthetic and 7 Design-Bench tasks. On synthetic tasks the surrogate main effect dominates ($\eta^2=0.37$ vs. 0.01 for the optimizer) with a large surrogate $\times$ optimizer interaction ($\eta^2=0.17$): gradient ascent collapses on the ensemble but is inert on the GP. Controls locate the cause in the GP’s posterior *mean*, not its calibration—the GP-ensemble gap is unchanged when the pessimism term is removed ($\beta=0$) and *widens* when the ensemble is given the GP’s data subsample. The ensemble’s LCB premise, which we show equals the bound’s validity, is moderately covered in-distribution (0.73, below the nominal 0.90) but collapses on the designs its own gradient ascent returns (0.41)—while remaining well covered (0.97) on the GP’s proposals, so the failure is the ensemble $\times$ gradient interaction, not a surrogate defect. Split-conformal recalibration restores in-distribution coverage to its 0.90 target but not the shifted proposal coverage; we package this as a coverage diagnostic. Finally, method choice is highly significant on synthetic benchmarks (Friedman $p=6 \times 10^{-5}$) but unresolved on Design-Bench ($p=0.69$; an equivalence test is underpowered at $N=7$): the differences these benchmarks reward do not transfer to real tasks.

1 Introduction

Offline model-based optimization (MBO) seeks a design x maximizing an expensive black-box f from a fixed dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$, with no further queries to f during search (Trabucco et al. 2022; Kim et al. 2025). The dominant recipe fits a surrogate $\hat{f}$ and optimizes it; because optimization pushes designs into regions where $\hat{f}$ is unconstrained, the surrogate is typically made conservative—by a learned regularizer (Trabucco et al. 2021) or by a pessimistic acquisition such as the lower confidence bound (LCB) $\mu(x) - \beta\sigma(x)$

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

(Srinivas et al. 2012). A common modeling intuition is that Gaussian-process LCB (GP-LCB) is preferable to neural-ensemble LCB at low dimension, on the grounds that the GP posterior is better calibrated.

That attribution is not identified. A GP-LCB pipeline and an ensemble-LCB pipeline differ in *two* coupled factors: the surrogate class (GP posterior vs. ensemble disagreement) and the acquisition optimizer (GPs are usually optimized by perturbation or quasi-Newton restarts, neural surrogates by gradient ascent on x). Existing systematic studies vary one axis with the other fixed: surrogate-class comparisons hold the optimizer constant (Tan et al. 2025)—a Bayesian-optimization study of GP vs. neural surrogates similarly fixes the acquisition and finds their ranking problem-dependent, with deep ensembles often the weakest (Li, Rudner, and Wilson 2024)—and the conservative-surrogate lineage holds gradient ascent fixed while varying the regularizer (Trabucco et al. 2021). The subfield’s own survey names realistic benchmarking and surrogate-uncertainty estimation among its open problems (Kim et al. 2025), yet whether reported gains come from the surrogate, the optimizer, or neither has not been measured under control. We supply that missing measurement.

Contributions.

1. **A controlled surrogate $\times$ optimizer decomposition** (Sections 3–4) over {ensemble, exact GP, sparse GP} $\times$ {gradient, perturbation, CMA-ES} on 7 synthetic and 7 Design-Bench tasks under one score-closure protocol. The surrogate main effect dominates on synthetic tasks ($\eta^2=0.37$ vs. 0.01) with a large surrogate $\times$ optimizer *interaction* ($\eta^2=0.17$); a matched-tuning control shows the surrogate effect is not a GP-tuning artifact (it retains 76%).
2. **The mechanism is inductive bias, not calibration** (Section 5). Three controls locate the GP’s advantage in its posterior *mean*: the GP-ensemble gap is unchanged with pessimism off ($\beta=0$), *widens* when the ensemble is given the GP’s data subsample, and the ensemble’s LCB premise—which we show (Prop. 1) equals the bound’s validity—collapses only on the ensemble’s own gradient proposals (coverage 0.41) while holding on the GP’s (0.97). We package the coverage measurement as a diagnostic with a split-conformal repair (Prop. 2).
3. **A benchmark-validity result** (Section 4), under a con-

trolled random-forest-oracle protocol. Method choice is highly significant on synthetic benchmarks (Friedman $p=6 \times 10^{-5}$) but statistically unresolved on Design-Bench ($p=0.69$; a paired equivalence test is underpowered at $N=7$ tasks): the method rankings synthetic suites reward do not transfer to the real tasks—a gap we show is not an artifact of the random-forest oracles. To our knowledge this is the first controlled measurement disentangling the surrogate and optimizer contributions in offline MBO, rather than a new optimizer or regularizer.

These threads share one cause. Contributions 2 and 3 are two views of a single prior–task match: the synthetic surfaces are globally smooth, matching the GP’s Matérn prior, so the GP wins decisively; the Design-Bench oracles are not, the match breaks, and both the GP edge and the method separation vanish.

2 Background and Related Work

LCB and the pessimism premise. GP-LCB $\mu - \beta\sigma$ has regret guarantees under GP assumptions (Srinivas et al. 2012); deep ensembles (Lakshminarayanan, Pritzel, and Blundell 2017) provide a cheap uncertainty proxy but no calibration guarantee. The pessimism principle underlying conservative offline methods is that a valid lower bound prevents the optimizer from exploiting model error (Jin, Yang, and Wang 2021; Rashidinejad et al. 2021; Xie et al. 2021). We make the premise behind LCB operationally measurable.

Offline MBO. Conservative Objective Models (Trabucco et al. 2021) add a CQL-style regularizer (Kumar et al. 2020); related methods include MINs (Kumar and Levine 2020), CbAS (Brookes, Park, and Listgarten 2019), autofocused oracles (Fannjiang and Listgarten 2020), DDOM (Krishnamoorthy, Mashkaria, and Grover 2023a), RoMA (Yu et al. 2021), BDI (Chen et al. 2022), BONET (Krishnamoorthy, Mashkaria, and Grover 2023b), and ranking surrogates (Tan et al. 2025). Optimizer-as-contribution work (Chemingui et al. 2024) argues search strategy is under-explored but proposes a method rather than decomposing. Design-Bench (Trabucco et al. 2022) standardizes tasks and, by fixing a protocol, documents the confounded status quo we dissect.

Calibration and conformal prediction. Post-hoc calibration (González et al. 2016) and conformal methods give distribution-free coverage (Angelopoulos and Bates 2023); conformal prediction sets have been used inside the Bayesian-optimization loop itself (Stanton, Maddox, and Wilson 2023). We use split-conformal and its weighted extension (Tibshirani et al. 2019) to repair the LCB premise and to diagnose where the repair transfers, positioning our calibration arm as *decompose and diagnose*: measuring coverage where the optimizer operates, rather than certifying a fixed pipeline.

Evaluation critiques and controlled decompositions. A recurring lesson from adjacent subfields is that reported gains can be artifacts of evaluation rather than method: reproducibility and statistical audits of deep RL (Henderson et al. 2018; Agarwal et al. 2021) and the demonstration that aggregate benchmark scores are structurally unreliable (Balduzzi

et al. 2018) each pair a critique with a reusable tool (a protocol, a stratified-bootstrap library, Nash averaging). Within offline optimization, SOO-Bench (Zhu, Qian et al. 2025) pursues a complementary benchmark-validity axis, stress-testing optimizer *stability* across problem instances. Our study is in this tradition—a controlled decomposition plus a coverage diagnostic—but targets the surrogate $\times$ optimizer confound in offline MBO, where it is not merely statistical noise but a specific, measurable premise (LCB coverage) that prior work has not quantified in the region the optimizer actually reaches.

3 The Decomposition Grid

Surrogates. (i) **Ensemble:** $K=5$ two-hidden-layer MLPs (width 96, ReLU), trained by MSE for 35 epochs (Adam, lr 3×10^{-3}); μ, σ are the member mean and standard deviation. (ii) **Exact GP:** a differentiable single-task GP (ARD Matérn-5/2), fit by marginal likelihood on a score-biased subsample ($N_{\max}=800$). (iii) **Sparse variational GP (SVGP):** 128 inducing points, ARD Matérn-5/2, 250 ELBO steps ($N_{\max}=2000$). The GPs additionally fit kernel hyperparameters by marginal likelihood—per-run tuning the ensemble does not receive; the matched-tuning control (below) freezes it to isolate the surrogate-class effect.

Optimizers. (i) **Gradient:** 100 Adam steps on x (lr 0.05), box-clipped to $[0, 1]^d$. (ii) **Perturbation:** 5 rounds of Gaussian hill-climbing ($\sigma \in \{0.1, 0.05, 0.02\}$). (iii) **CMA-ES:** population search (initial $\sigma=0.2$, separable variant when $d>500$). All three surrogates are differentiable (the exact GP is a differentiable single-task GP), so every surrogate $\times$ optimizer combination is realizable: a full 3×3 grid of 9 cells, plus the COMs, CbAS, and gradient-ascent baselines.

Acquisition. All cells maximize the LCB $\mu(x) - \beta\sigma(x)$ with $\beta=2$; $\beta=0$ recovers pure surrogate maximization.

Score-closure protocol. Every cell and baseline shares one evaluation closure: fit the surrogate on $\mathcal{D}$, form the LCB acquisition, hand it to the optimizer, collect the 128 proposed designs, and score them with the ground-truth oracle. Only the two swept factors—the surrogate producing (μ, σ) and the optimizer searching the acquisition—vary across cells; the data split, candidate budget, input normalization, and oracle scoring are held identical. This shared closure is what licenses attributing score differences to the surrogate $\times$ optimizer factors rather than to incidental protocol choices, and it is precisely the control that pipeline-level comparisons of full published methods lack.

Tasks and protocol. 7 synthetic functions (Branin-2D through Griewank-30D; fixed dataset drawn once at seed 0) and 7 Design-Bench tasks (TF-Bind-8/10 with exact oracles; Superconductor, GFP, UTR, Ant, D’Kitty with random-forest oracles to remove simulator/framework dependencies). Discrete designs are relaxed to per-position class logits and decoded by argmax; continuous inputs and all scores are min-max normalized. Each method proposes 128 candidates;

we report the 100th-percentile (max) oracle score, following Design-Bench. We run 30 seeds on synthetic and 16 on Design-Bench.

Matched tuning (identifiability control). Exact and sparse GPs, unlike ensembles, can spend per-run hyperparameter-optimization budget. The *matched* arm freezes GP hyperparameter fitting so every surrogate gets the same zero per-run tuning budget, isolating the surrogate-class effect from a GP-tuning advantage.

Attribution. Per task we min-max normalize the grid and compute a two-way ANOVA: η_{surr}^2 and η_{opt}^2 are the fractions of normalized-score variance explained by the surrogate and optimizer main effects. Significance uses Wilcoxon signed-rank with Holm correction, an omnibus Friedman test, Nemenyi critical differences, and bootstrap rank confidence intervals.

4 Results: What Drives Performance

The optimizer effect is an ensemble-specific interaction. Table 1 shows the mechanism directly. Within the *ensemble* rows, changing the optimizer moves the score by orders of magnitude (Branin -0.78 for perturbation vs. -9.27 for gradient; Griewank -395 vs. -2592): gradient ascent drives designs into regions where the ensemble over-estimates, collapsing to the domain boundary. Within the *GP* rows the same optimizer switch barely moves the score (Branin all ≈ -0.40). Figure 2 summarizes this as the optimizer-induced spread per surrogate.

Attribution: surrogate main effect, plus a large interaction. The two-way ANOVA assigns the synthetic-task variance to the surrogate main effect ($\eta_{\text{surr}}^2=0.37$ vs. $\eta_{\text{opt}}^2=0.01$): averaging over optimizers, GP and SVGP lead the ensemble by a real marginal-mean gap. Under *matched* tuning the surrogate effect persists ($\eta_{\text{surr}}^2=0.28$, retaining 76% of its unmatched 0.37), so it is not merely a per-run tuning budget the ensemble lacks. The optimizer main effect is small because its influence is concentrated on one surrogate: the surrogate $\times$ optimizer *interaction* is $\eta_{\text{inter}}^2=0.17$, an order of magnitude above the optimizer main effect (Table 2). This is the ensemble-specific collapse quantified—the optimizer matters, but only through the ensemble.

Significance, and the synthetic $\rightarrow$ real collapse. On synthetic tasks the nine grid cells are strongly separated: a Friedman omnibus over the 9 cells and 7 tasks gives $p=6.1 \times 10^{-5}$, with the best cell (GP $\times$ gradient) at mean rank 2.3 (bootstrap 95% CI [1.3, 3.6]) against a Nemenyi critical difference of 4.5. On Design-Bench the same analysis yields $p=0.69$: no cell is distinguishable from any other (best cell GP $\times$ perturbation at mean rank 3.6, CI [1.9, 5.7]; the whole grid lies within one critical difference). Figure 3 contrasts the two regimes, and the two-panel decomposition map (Figure 1) shows the same collapse at a glance: the sharp dark/bright rank structure on the left flattens to a uniform mid-gray on the right. The driver also reverses: on real tasks the (small) surrogate advantage shrinks ($\eta_{\text{surr}}^2=0.05$) and, if anything, the optimizer explains marginally more

($\eta_{\text{opt}}^2=0.08$)—no factor explains much variance (Table 2). We state this as a benchmark-validity finding: the large, significant method differences synthetic suites reward are *statistically unresolved* on the real tasks in Design-Bench. We are careful not to overstate the null: a paired equivalence test (TOST) on the best-versus-worst cells does *not* establish equivalence—the 90% CI on their gap is ± 0.48 normalized units—so the real-task result is that method choice is unresolved and *underpowered at $N=7$ tasks*, not that the methods are provably equal. Table 3 gives the full Design-Bench grid. Two caveats sharpen it: the ensemble $\times$ gradient collapse is a continuous-domain phenomenon—on the two exact-oracle tasks (TF-Bind-8/10) that cell does not collapse but ties or leads (e.g. 2.20 on TF-Bind-8, above every GP cell)—and elsewhere the methods cluster in a narrow band with the per-task best split across surrogates and optimizers.

5 Mechanism: Inductive Bias, Not Calibration

Why does the GP win, and why does gradient ascent collapse on the ensemble but not on the GP? The intuitive answer—that the GP’s uncertainty is better calibrated, so its pessimistic bound is valid—is wrong. Three controls localize the effect to the GP’s posterior *mean*, an inductive bias, and to an ensemble $\times$ gradient interaction; σ -calibration plays no causal role.

The advantage survives with pessimism off. If the GP’s edge were a calibrated LCB, removing the pessimism term should erase it. It does not. Re-running the full grid at $\beta=0$ (posterior-mean maximization, no σ penalty), the GP–ensemble marginal gap is essentially unchanged—0.51 at $\beta=2$ versus 0.47 at $\beta=0$ (task-normalized, averaged over optimizers; the paired difference has 95% CI $[-0.02, 0.10]$, indistinguishable from zero; Table 4)—and the ensemble still collapses under gradient ascent with no pessimism at all (Griewank -2701 vs. the GP’s -0.94). The advantage is a property of the GP’s smooth mean, not of σ .

The advantage is not the GP’s data subsample. The GP fits a score-biased 800-point subsample while the ensemble uses all data. Giving the ensemble the *same* subsample does not close the gap—it *widens* it, from 0.51 to 0.76: the ensemble does worse with fewer points and still collapses under gradient. The GP wins *despite* fitting less data, isolating the effect to the surrogate class.

The ensemble fails only under its own gradient ascent. The collapse is not a blanket ensemble defect. The pessimism premise $\{\mu - f \leq \beta\sigma\}$ (Prop. 1) is moderately covered for the ensemble in-distribution (0.73, below the nominal 0.90) but collapses on the designs its *own* gradient ascent returns (0.41 mean; 0.00 on the collapse tasks). On the GP’s proposals, however, the ensemble premise is well covered (0.97): the ensemble is not uniformly miscalibrated—it fails specifically where its gradient optimizer drives it (Figure 4). The GP’s premise holds both in-distribution (0.98) and on its own proposals (0.97): its smooth mean offers no over-estimated argmax for gradient ascent to exploit. This is the

Table 1: Offline MBO on synthetic tasks: 100th-percentile score by surrogate $\times$ optimizer (30 seeds; higher is better). **Bold** = best per task. Within the ensemble rows the optimizer swings the score enormously (Branin $-0.78 \rightarrow -9.27$; Griewank $-395 \rightarrow -2612$); within the GP rows it is nearly inert (Branin all ≈ -0.40). This interaction, not a surrogate-only or optimizer-only main effect, is the source of the reported GP-LCB “advantage.”

Surrogate $\times$ Opt.	Branin	Styblinski	Levy	Rosenbrock	Rastrigin	Ackley	Griewank
Ens $\times$ Grad	-9.27	6.37	-2.14	-0.28	-7.71	-3.66	-2592
Ens $\times$ Pert	-0.78	33.08	-0.40	-0.12	-8.44	-6.32	-395
Ens $\times$ CMA	-14.01	5.21	-3.19	-0.48	-10.81	-4.31	-2613
GP $\times$ Grad	-0.40	27.57	-0.05	-0.09	-4.83	-0.55	-0.94
GP $\times$ Pert	-0.40	36.15	-0.24	-0.08	-8.28	-6.28	-269
GP $\times$ CMA	-0.40	26.65	-0.05	-0.09	-5.06	-0.59	-1.00
SVGP $\times$ Grad	-0.45	11.83	-0.08	-0.04	-2.83	-0.69	-2.11
SVGP $\times$ Pert	-0.40	34.35	-0.25	-0.08	-8.27	-6.14	-275
SVGP $\times$ CMA	-0.53	11.60	-0.08	-0.04	-3.00	-0.73	-2.17

Table 2: Two-way ANOVA attribution (η^2 , fraction of task-normalized variance explained by each term) across regimes. The surrogate main effect dominates on synthetic tasks and survives the matched-tuning control; the surrogate $\times$ optimizer interaction is itself an order of magnitude above the optimizer main effect. On Design-Bench neither factor explains much variance and the (small) ordering reverses. Task-and-seed bootstrap 95% CIs (synthetic, unmatched): $\eta_{\text{surr}}^2 \in [0.25, 0.57]$, $\eta_{\text{opt}}^2 \in [0.01, 0.19]$ (non-overlapping), $\eta_{\text{inter}}^2 \in [0.11, 0.26]$.

Regime	η_{surr}^2	η_{opt}^2	η_{inter}^2	Friedman p
Synthetic (unmatched)	0.37	0.01	0.17	6.1×10^{-5}
Synthetic (matched tuning)	0.28	0.02	0.12	—
Design-Bench	0.05	0.08	0.01	0.69

interaction the grid isolates, now measured directly rather than inferred.

Coverage is validity. These frequencies are not incidental; they exactly measure the validity of the pessimistic bound. Write $L_\beta(x) = \mu(x) - \beta\sigma(x)$; call L_β a valid $(1-\delta)$ lower bound under a distribution Q if $\Pr_{x \sim Q}(f(x) \geq L_\beta(x)) \geq 1-\delta$.

Proposition 1 (Coverage of the premise is LCB validity). *For any Q, β, δ , $\Pr_{x \sim Q}(f(x) \geq L_\beta(x)) = \Pr_{x \sim Q}(\mu(x) - f(x) \leq \beta\sigma(x))$. Hence L_β is a valid $(1-\delta)$ lower bound under Q iff the premise $\{\mu - f \leq \beta\sigma\}$ has Q -probability at least $1-\delta$.*

The one-line identity ($\sigma > 0$; the two events coincide) turns “the bound holds” into a measurable frequency, evaluated under the data $\mathcal{D}$ and the proposal Π . At $\beta=2$, ensemble coverage is 0.73 in-distribution / 0.41 on proposals (synthetic) and 0.77 / 0.18 (real), below the nominal 0.90 exactly on the proposals where the bound must hold (Figure 5; per-task values in the supplement).

Pessimism is regularization, not calibration. Two further observations confirm calibration is not the driver. The uncertainty σ is a weak error signal—Spearman ρ between σ and $|\mu - f|$ is only ≈ 0.1 —yet increasing β improves the score

on 6 of 7 synthetic tasks (Figure 6; median normalized slope $+0.19$). Pessimism helps as crude distance-regularization, penalizing high-variance far-from-data regions, not as calibrated conservatism; the lone exception (Ackley) is where the penalty most often points the wrong way. Per-task calibration quality does not predict whether pessimism helps. Ensemble size is the standard $K=5$ (Lakshminarayanan, Pritzel, and Blundell 2017); shrinking K inflates σ and the pessimism penalty, improving the ensemble on its own K -sweep—corroborating the regularization reading, orthogonal to the $\beta=0$ result.

A distribution-free repair, and where it transfers. Split-conformal calibration replaces the arbitrary β with a data-driven multiplier.

Proposition 2 (Conformal repair; shift-limited transfer). *Let $\{(x_i, f(x_i))\}_{i=1}^n$ be exchangeable from P , $r_i = (\mu(x_i) - f(x_i)) / \sigma(x_i)$ the signed one-sided nonconformity, and $\hat{q}$ the $[(n+1)(1-\delta)]$ -th smallest r_i . Then for a fresh $x \sim P$, $\Pr(f(x) \geq \mu(x) - \hat{q}\sigma(x)) \geq 1-\delta$. Under a shifted proposal $\Pi \neq P$ with density ratio $w = d\Pi/dP$, validity is restored by weighting the calibration quantile by w (weighted conformal; Tibshirani et al. 2019).*

Empirically the fitted multiplier varies widely across tasks (synthetic $\hat{q} \in [1.8, 16]$), quantifying how far raw σ undercovers. Replacing $\beta\sigma$ by $\hat{q}\sigma$ restores in-distribution coverage to its 0.90 target on every task (mean 0.90 synthetic and real), exactly as Prop. 2 predicts; but on the shifted proposal Π coverage stays low (0.51 / 0.31 mean, near zero on the sharpest-shift tasks), because $\Pi \neq P$. This yields a concrete recipe (Algorithm 1): a *coverage diagnostic* that reports in-distribution and proposal coverage, flags an unreliable LCB when they diverge, and a *weighted-conformal repair* when a density-ratio estimate is available. The negative transfer is itself the actionable signal: an LCB whose proposal coverage is near zero should not be trusted regardless of β .

6 Discussion, Guidance, and Limitations

Practitioner guidance. (i) On smooth, low-to-moderate-dimensional problems, a GP/SVGP surrogate is robust to the acquisition optimizer (its smooth mean offers no exploitable argmax); a neural ensemble is not, and should be paired with

Surrogate class dominates on synthetic tasks; the grid flattens on Design-Bench




Figure 1: Decomposition map: mean rank (over tasks, 1=best) of every surrogate×optimizer cell. **Left (synthetic):** the GP / SVGP surrogate is best regardless of optimizer (dark cells), while the ensemble is poor and swings with the optimizer (bright cells)—the surrogate main effect plus the ensemble×optimizer interaction. **Right (Design-Bench):** the same grid collapses to a near-uniform mid-rank field, mirroring the statistically non-significant omnibus.




Figure 2: Optimizer-induced spread by surrogate class (synthetic; per-seed task-normalized 100th-percentile score, higher is better; boxes span the interquartile range). The three optimizers produce sharply different distributions on the ensemble surrogate but nearly identical, tightly-clustered distributions on the GP and SVGP surrogates: the acquisition optimizer matters enormously for ensembles and negligibly for the GP and SVGP surrogates, whose smooth posterior mean (not calibration) is the cause (Section 5).

a conservative optimizer (perturbation/CMA-ES) rather than aggressive gradient ascent. (ii) Before trusting an LCB, run the coverage diagnostic (Algorithm 1); if proposal coverage is near zero, the pessimism is not doing what the theory says, and neither β -tuning nor in-distribution conformal calibration fixes it. (iii) On real Design-Bench tasks, method choice within this family is not statistically resolvable at standard




Figure 3: Critical-difference diagram (Nemenyi, $\alpha=0.05$) over the 9 grid cells. Synthetic tasks separate the cells sharply (Friedman $p=6.1 \times 10^{-5}$); Design-Bench does not ($p=0.69$), all cells falling within one critical difference. Method rankings do not transfer from synthetic to real.

seed counts—effort is better spent on the surrogate’s informativeness than on the search rule.

Anticipated objections. Three objections deserve direct answers. (i) *That these benchmarks are imperfect is already known.* The contribution is its first controlled measurement—a crossed surrogate×optimizer decomposition with a matched-tuning control, and a premise-coverage frequency not previously reported. (ii) *There is no new state*

Table 3: Offline MBO on Design-Bench (16 seeds; normalized 100th-percentile score, higher is better; random-forest oracles for the non-exact tasks). **Bold** = best per task including the COMs/CbAS baselines. Unlike the synthetic grid, the methods cluster tightly and the per-task winner is split across surrogates and optimizers (Friedman $p=0.69$ over the grid); the ensemble/COMs gradient-ascent collapse persists only on GFP.

Surrogate×Opt.	TF-Bind-8	TF-Bind-10	Superconductor	GFP	UTR	Ant	D’Kitty
Ens × Grad	2.20	1.34	0.99	−9.61	1.01	0.92	0.74
Ens × Pert	1.00	1.31	1.29	2.48	0.98	1.29	1.05
Ens × CMA	2.06	1.15	1.18	−5.63	0.99	0.89	0.72
GP × Grad	1.00	1.31	1.24	2.45	0.99	1.52	1.06
GP × Pert	1.00	1.31	1.31	2.45	0.99	1.52	1.11
GP × CMA	1.00	1.00	1.26	2.46	0.99	1.52	1.02
SVGP × Grad	1.84	1.11	1.07	1.80	0.99	0.99	1.02
SVGP × Pert	1.45	1.31	1.25	2.45	0.99	1.52	1.09
SVGP × CMA	1.83	1.11	1.26	2.48	0.93	0.91	1.02
COMs (baseline)	2.21	1.35	1.01	−9.20	1.01	0.98	0.95
CbAS (baseline)	2.12	1.28	1.36	1.85	0.99	1.53	1.02

Table 4: Gap controls (synthetic; task-normalized GP − ensemble marginal, with 10-seed task+seed bootstrap 95% CIs). The gap is unchanged with pessimism off ($\beta=0$; the paired $\beta=2$ vs. $\beta=0$ difference has CI $[-0.02, 0.10]$) and widens when the ensemble is given the GP’s score-biased subsample: the edge is the surrogate’s posterior mean, not σ -calibration or data.

Setting	gap	95% CI
$\beta=2$ (pessimism on)	0.51	[0.43, 0.58]
$\beta=0$ (pessimism off)	0.47	[0.37, 0.57]
ensemble on the GP’s 800-pt subsample	0.76	[0.29, 1.32]

Algorithm 1 Coverage diagnostic and conformal-LCB repair

Input: surrogate (μ, σ) , data $\mathcal{D}$, level δ , proposals Π

Output: multiplier $\hat{q}$, coverages $(\hat{c}_{\text{in}}, \hat{c}_{\text{ood}})$, flag

- 1: Split $\mathcal{D}$ into fit / calibration folds.
- 2: $r_i \leftarrow (\mu(x_i) - f(x_i)) / \sigma(x_i)$ on calibration fold (signed, one-sided; floor σ).
- 3: $\hat{q} \leftarrow \lceil (n+1)(1-\delta) \rceil$ -th smallest r_i .
- 4: $\hat{c}_{\text{in}} \leftarrow \widehat{\Pr}_{x \sim \mathcal{D}}(f \geq \mu - \hat{q}\sigma)$.
- 5: $\hat{c}_{\text{ood}} \leftarrow \widehat{\Pr}_{x \sim \Pi}(f \geq \mu - \hat{q}\sigma)$.
- 6: **if** $\hat{c}_{\text{ood}} < 1-\delta$ **then**
- 7: **flag** LCB unreliable; if $w=d\Pi/dP$ available, reweight r_i by w and refit $\hat{q}$.
- 8: **end if**
- 9: **return** $\hat{q}, (\hat{c}_{\text{in}}, \hat{c}_{\text{ood}})$, flag

of the art. The constructive artifact is the coverage diagnostic and weighted-conformal repair (Algorithm 1): a reusable check, the same form of contribution as the evaluation tools this line of work is built on. (iii) *The real-task result is null.* It is a significance collapse, not a claim of equivalence—the differences are unresolved, not provably zero (underpowered at $N=7$)—a benchmark-validity statement that the rewarded differences do not transfer.




Figure 4: Premise coverage $\Pr(f \geq \mu - \beta\sigma)$ at $\beta=2$ (synthetic mean), each surrogate on in-distribution points, its own proposals, and the other surrogate’s proposals. The ensemble premise collapses (0.41) *only* on the designs its own gradient ascent returns; on the GP’s proposals it is well covered (0.97), and the GP premise holds everywhere. The failure is the ensemble×gradient interaction, not a surrogate defect.

Scope and external validity. The synthetic→real collapse is a statement about the *measured* regime. Two boundary conditions bound it. First, for the non-exact tasks we substitute a random-forest oracle to remove simulator and framework dependence—a genuine substitution on GFP, UTR, Ant, and D’Kitty (the native oracle is an RF only for Superconductor), so a different oracle could re-separate methods. But it does not manufacture the null: on the three tasks with no smoothing RF substitution (TF-Bind-8/10 exact, Superconductor native RF), the omnibus stays flat (Friedman $p=0.93$) and the median 9-cell spread (0.34) is no smaller than on the substituted tasks (0.39). Second, the real-task omnibus is over 7 tasks, and the equivalence test is underpowered there (Section 4): we can say the method differences are

The pessimism premise fails on proposed designs; conformal repairs in-distribution, not OOD




Figure 5: Coverage of the pessimism premise $\mu - f \leq \beta\sigma$ vs. β on synthetic (left) and Design-Bench (right) tasks: in-distribution (solid) and on the optimizer’s OOD proposals (dashed), with ± 1 s.d. bands across tasks and the split-conformal repair shown as stars (*). Ensemble in-distribution coverage is moderate, but proposal (OOD) coverage is far lower and task-dependent—near zero on the tasks where gradient ascent collapses; conformal recalibration lifts in-distribution coverage to its 0.90 target (upper star) but leaves the OOD coverage the bound actually needs far below it (lower star).




Figure 6: Pessimism sweep: per-task normalized 100th-percentile score vs. β (faint lines, individual synthetic tasks; bold line, mean). Score rises with β on 6 of 7 tasks even though σ is an uninformative error signal ($\rho \approx 0.1$)—pessimism acts as distance-regularization. Ackley is the single task where the penalty points the wrong way.

unresolved, not that they are zero.

Limitations. Our real-task evidence is Design-Bench with random-forest oracles for the non-exact tasks; the null is scoped to that suite and to CPU-scale training. A cross-check ran official COMs and CbAS: our CbAS matches the official scores ($|\Delta| = 0.004$ on TF-Bind-8), while our COMs and official COMs diverge ($|\Delta| = 1.2$, a reduced-epoch official run and a different oracle variant), so we report both side by side. The synthetic datasets are fixed across seeds (seed-0), so reported variance is training/optimization, not data-draw. Premise coverage on the discrete tasks uses relaxed logits,

not decoded sequences, so those numbers (notably GFP) are softer—we lean the coverage mechanism on the synthetic tasks. Weighted conformal needs a density-ratio estimate we do not learn; we characterize the transfer gap rather than close it.

7 Conclusion

The reported low-dimensional advantage of GP-LCB in offline MBO is not a clean surrogate-class effect and not a calibration effect: it is a surrogate $\times$ optimizer interaction driven by the GP’s smooth posterior mean, on which gradient ascent finds no over-estimated argmax to exploit as it does on the ensemble’s jagged mean. Removing pessimism ($\beta=0$) leaves the edge intact and giving the ensemble the GP’s data does not close it—the cause is inductive bias, not uncertainty. Making the premise measurable (coverage equals LCB validity) turns it into a diagnostic: a bound whose proposal coverage is near zero should not be trusted regardless of β . And the differences these synthetic benchmarks reward are unresolved on real tasks—the same prior-task match, seen from both sides. We contribute the decomposition, the mechanism and its controls, the coverage diagnostic, and the synthetic-to-real validity result.

Ethical Statement

This work studies optimization methodology on public benchmarks and introduces no new data or deployed system. By flagging when a pessimistic bound is invalid, the coverage diagnostic aims to reduce over-confident designs downstream.

References

- Agarwal, R.; Schwarzer, M.; Castro, P. S.; Courville, A.; and Bellemare, M. G. 2021. Deep Reinforcement Learning at the Edge of the Statistical Precipice. In *Advances in Neural Information Processing Systems*.
- Angelopoulos, A. N.; and Bates, S. 2023. Conformal Prediction: A Gentle Introduction. *Foundations and Trends in Machine Learning*, 16(4): 494–591.
- Balduzzi, D.; Tuyls, K.; Perolat, J.; and Graepel, T. 2018. Re-evaluating Evaluation. In *Advances in Neural Information Processing Systems*.
- Brookes, D. H.; Park, H.; and Listgarten, J. 2019. Conditioning by Adaptive Sampling for Robust Design. In *International Conference on Machine Learning*, 773–782. PMLR.
- Chemingui, Y.; Deshwal, A.; Hoang, T. N.; and Doppa, J. R. 2024. Offline Model-Based Optimization via Policy-Guided Gradient Search. In *Proceedings of the AAAI Conference on Artificial Intelligence*.
- Chen, C.; Zhang, Y.; Fu, J.; Li, X.; and Poupart, P. 2022. Bidirectional Learning for Offline Infinite-width Model-Based Optimization. In *Advances in Neural Information Processing Systems*, volume 35, 6800–6814.
- Fannjiang, C.; and Listgarten, J. 2020. Autofocused Oracles for Model-Based Design. In *Advances in Neural Information Processing Systems*, volume 33, 12945–12956.
- González, J.; Dai, Z.; Hennig, P.; and Lawrence, N. D. 2016. Batch Bayesian Optimization via Local Penalization. In *International Conference on Artificial Intelligence and Statistics*, 648–657. PMLR.
- Henderson, P.; Islam, R.; Bachman, P.; Pineau, J.; Precup, D.; and Meger, D. 2018. Deep Reinforcement Learning that Matters. In *Proceedings of the AAAI Conference on Artificial Intelligence*.
- Jin, Y.; Yang, Z.; and Wang, Z. 2021. Is Pessimism Provably Efficient for Offline Reinforcement Learning? In *International Conference on Machine Learning*, 5084–5096. PMLR.
- Kim, M.; Gu, J.; Yuan, Y.; Yun, T.; Liu, Z.; Bengio, Y.; and Chen, C. 2025. Offline Model-Based Optimization: Comprehensive Review. *arXiv preprint arXiv:2503.17286*.
- Krishnamoorthy, S.; Mashkaria, S. M.; and Grover, A. 2023a. Diffusion Models for Black-Box Optimization. In *International Conference on Machine Learning*, 17842–17857. PMLR.
- Krishnamoorthy, S.; Mashkaria, S. M.; and Grover, A. 2023b. Generative Pretraining for Black-Box Optimization. *arXiv preprint arXiv:2206.10786*.
- Kumar, A.; and Levine, S. 2020. Model Inversion Networks for Model-Based Optimization. In *Advances in Neural Information Processing Systems*, volume 33, 5126–5137.
- Kumar, A.; Zhou, A.; Tucker, G.; and Levine, S. 2020. Conservative Q-Learning for Offline Reinforcement Learning. In *Advances in Neural Information Processing Systems*, volume 33, 1179–1191.
- Lakshminarayanan, B.; Pritzel, A.; and Blundell, C. 2017. Simple and Scalable Predictive Uncertainty Estimation using Deep Ensembles. In *Advances in Neural Information Processing Systems*, volume 30.
- Li, Y. L.; Rudner, T. G. J.; and Wilson, A. G. 2024. A Study of Bayesian Neural Network Surrogates for Bayesian Optimization. In *International Conference on Learning Representations*.
- Rashidinejad, P.; Zhu, B.; Ma, C.; Jiao, J.; and Russell, S. 2021. Bridging Offline Reinforcement Learning and Imitation Learning: A Tale of Pessimism. In *Advances in Neural Information Processing Systems*, volume 34, 11702–11716.
- Srinivas, N.; Krause, A.; Kakade, S. M.; and Seeger, M. W. 2012. Gaussian Process Optimization in the Bandit Setting: No Regret and Experimental Design. *IEEE Transactions on Information Theory*, 58(5): 3250–3265.
- Stanton, S.; Maddox, W.; and Wilson, A. G. 2023. Bayesian Optimization with Conformal Prediction Sets. In *International Conference on Artificial Intelligence and Statistics*. PMLR.
- Tan, R.-X.; Xue, K.; Lyu, S.-H.; Shang, H.; Wang, Y.; Wang, Y.; Fu, S.; and Qian, C. 2025. Offline Model-Based Optimization by Learning to Rank. In *International Conference on Learning Representations*.
- Tibshirani, R. J.; Foygel Barber, R.; Candes, E. J.; and Ramdas, A. 2019. Conformal Prediction Under Covariate Shift. In *Advances in Neural Information Processing Systems*.
- Trabucco, B.; Geng, X.; Kumar, A.; and Levine, S. 2022. Design-Bench: Benchmarks for Data-Driven Offline Model-Based Optimization. In *International Conference on Machine Learning*, 21658–21676. PMLR.
- Trabucco, B.; Kumar, A.; Geng, X.; and Levine, S. 2021. Conservative Objective Models for Effective Offline Model-Based Optimization. In *International Conference on Machine Learning*, 10358–10368. PMLR.
- Xie, T.; Cheng, C.-A.; Jiang, N.; Mineiro, P.; and Agarwal, A. 2021. Bellman-Consistent Pessimism for Offline Reinforcement Learning. In *Advances in Neural Information Processing Systems*, volume 34, 6683–6694.
- Yu, S.; Ahn, S.; Song, L.; and Shin, J. 2021. RoMA: Robust Model Adaptation for Offline Model-Based Optimization. In *Advances in Neural Information Processing Systems*, volume 34, 4619–4631.
- Zhu, Y.; Qian, H.; et al. 2025. SOO-Bench: Benchmarks for Evaluating the Stability of Offline Black-Box Optimization. In *International Conference on Learning Representations*.