

Decomposing the GP Advantage in Offline MBO

Anonymous submission

Abstract

Benchmark audits that name unreported implementation choices and show the ranking move are an established genre, and we claim no credit for its shape. Unclaimed is the composition: five confounds specific to offline model-based optimization surrogate comparison, a removal protocol, and a two-way crossed surrogate $\times$ optimizer variance decomposition. Correcting the two protocol confounds, target scaling and the candidate/oracle rule, does not shrink the reported surrogate advantage but strengthens it, moving η_{surr}^2 from a published 0.37 to 0.405 [0.290, 0.556]; the ensemble-size confound moves it the other way, so the strengthening is the net of the two protocol confounds, not a property of de-confounding in general. The four corner intervals overlap and are not resolvable at $n=7$ tasks. The firmer result is a sensitivity: η_{surr}^2 rises from 0.184 [0.085, 0.354] at $\beta=0$ to 0.406 [0.285, 0.564] at $\beta=2$, and doubles again with search budget. The mechanism section is a set of eliminations. With pessimism removed the GP family retains 61% of its advantage (0.319 [0.196, 0.460] against 0.525 [0.406, 0.614]), so pessimism amplifies a mean-quality base rather than creating it; sweeping member width $10.7\times$ at fixed K leaves the gap unchanged; the ensemble is the more accurate surrogate in-distribution, beating the GP’s held-out RMSE in 26 of 28 (task, width) cells, and still loses. A pre-registered bidirectional manipulation then refutes the account we had published: smoothing the ensemble’s mean by up to 98% does not close the gap but widens it, and drives premise coverage to 0.998 with no effect on the gap at all. Six candidate explanations are eliminated and none confirmed; what differs is which off-distribution maximizers a surrogate’s mean admits, not how rough it is. Under a matched query budget the optimizer main effect is $\eta_{\text{opt}}^2=0.038$ [0.003, 0.123], eight times the published value: established at high budget, underpowered at low. On Design-Bench we detect no difference at this power, which at $n=7$ tasks is not a demonstration of equivalence, partly because some surrogate cells are frozen and cannot move.

1 Introduction

Offline model-based optimization (MBO) seeks a design x maximizing an expensive black-box f from a fixed dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$, with no further queries to f during search

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

(Trabucco et al. 2022; Kim et al. 2025). The dominant recipe fits a surrogate $\hat{f}$ and optimizes it; because optimization pushes designs into regions where $\hat{f}$ is unconstrained, the surrogate is typically made conservative—by a learned regularizer (Trabucco et al. 2021) or by a pessimistic acquisition such as the lower confidence bound (LCB) $\mu(x) - \beta\sigma(x)$ (Srinivas et al. 2012). A common modeling intuition is that Gaussian-process LCB (GP-LCB) is preferable to neural-ensemble LCB at low dimension, on the grounds that the GP posterior is better calibrated.

That attribution is not identified. A GP-LCB pipeline and an ensemble-LCB pipeline differ in *two* coupled factors: the surrogate class (GP posterior vs. ensemble disagreement) and the acquisition optimizer (GPs are usually optimized by perturbation or quasi-Newton restarts, neural surrogates by gradient ascent on x). Existing systematic studies vary one axis with the other fixed: surrogate-class comparisons hold the optimizer constant (Tan et al. 2025)—a Bayesian-optimization study of GP vs. neural surrogates similarly fixes the acquisition and finds their ranking problem-dependent, with deep ensembles often the weakest (Li, Rudner, and Wilson 2024)—and the conservative-surrogate lineage holds gradient ascent fixed while varying the regularizer (Trabucco et al. 2021). The subfield’s own survey names realistic benchmarking and surrogate-uncertainty estimation among its open problems (Kim et al. 2025), yet whether reported gains come from the surrogate, the optimizer, or neither has not been measured under control. We supply that missing measurement.

What we concede. The move we make—name unreported implementation choices, give the protocol that removes them, show the ranking move—is an established genre, and we claim no credit for its shape. Its canonical instances are well known: reproducibility and statistical audits of deep RL (Henderson et al. 2018; Agarwal et al. 2021), the neural-recommender audit (Ferrari Dacrema, Cremonesi, and Jannach 2019), the metric-learning reality check (Musgrave, Belongie, and Lim 2020), and the large-scale GAN comparison (Lucic et al. 2018). One of them appeared at this venue. Conceding the shape is what makes the residual credible.

What is left. Unclaimed is the *composition*: a five-confound set specific to offline-MBO surrogate comparison, a protocol that removes them, and a two-way crossed

surrogate $\times$ optimizer variance decomposition. The one-way analogue is owned by fANOVA (Hutter, Hoos, and Leyton-Brown 2014); the closest two-axis design declines the two-way decomposition on the grounds that OFAT analysis cannot detect interactions (Moosbauer et al. 2022). The residual is one of composition, not of kind.

Contributions.

1. **Five named confounds and a removal protocol** (Section 3). Each is given a code-level trace and a signed effect on η_{surr}^2 .
2. **The net result runs against the genre’s direction** (Section 4). Correcting the two *protocol* confounds moves η_{surr}^2 from a published 0.37 to 0.405 [0.290, 0.556]—a strengthening, where audits normally shrink (Melis, Dyer, and Blunsom 2018); the ensemble-size confound moves it the other way, so the strengthening is the net of the two protocol confounds and not a property of de-confounding in general. We report that the four corner intervals overlap and are not resolvable at $n=7$ tasks, and lead instead with the β -dependence, which is.
3. **A mechanism narrowed by elimination** (Section 5), applied and diagnosed rather than discovered: distance-aware GP uncertainty (Liu et al. 2020) and the local-search reading of a UCB acquisition (Fan et al. 2024) are prior results we build on. Six controls remove σ , width, budget, held-out accuracy, mean roughness and premise coverage as explanations of the gap—the last two by a pre-registered manipulation that refutes our own prior account.
4. **A Design-Bench null, with its power stated** (Section 6). We detect no difference between methods at $n=7$ tasks—below the threshold for the omnibus test we run (Demšar 2006)—and report this as no detectable difference at this power, never as equivalence (Agarwal et al. 2021). Part of the reason is structural: some surrogate cells are frozen at a fixed dataset reference and cannot move at all.

What we do not claim. We make no field-reversal claim. That the statistical model often matters more than the acquisition heuristic is the organizing doctrine of the standard Bayesian-optimization review (Shahriari et al. 2016) and is a decade old; our optimizer-axis result falsifies the *local* premise of one recent paper (Chemingui et al. 2024) and nothing wider. We claim no mechanism discovery (Liu et al. 2020; Fan et al. 2024). We claim no equivalence on Design-Bench. And we present ensemble size K as a confound to be controlled, never as a discovered mechanism.

2 Background and Related Work

LCB and the pessimism premise. GP-LCB $\mu - \beta\sigma$ has regret guarantees under GP assumptions (Srinivas et al. 2012); deep ensembles (Lakshminarayanan, Pritzel, and Blundell 2017) provide a cheap uncertainty proxy but no calibration guarantee. The pessimism principle underlying conservative offline methods is that a valid lower bound prevents the optimizer from exploiting model error (Jin, Yang, and Wang 2021). We make that premise operationally measurable.

Offline MBO. The conservative-surrogate lineage runs from Conservative Objective Models (Trabucco et al. 2021), which add a CQL-style regularizer (Kumar et al. 2020), through generative and ranking-based surrogates (Brookes, Park, and Listgarten 2019; Krishnamoorthy, Mashkaria, and Grover 2023; Tan et al. 2025); all hold the search rule fixed while varying the model. Optimizer-as-contribution work (Chemingui et al. 2024) argues the converse—that search strategy is under-explored—but proposes a method rather than decomposing. Design-Bench (Trabucco et al. 2022) standardizes tasks and, by fixing a protocol, documents the confounded status quo we dissect.

Calibration and coverage. Conformal methods give distribution-free coverage, including inside the Bayesian-optimization loop (Angelopoulos and Bates 2023; Stanton, Maddox, and Wilson 2023). We use coverage of the LCB premise only as a *measurement*, taken where the optimizer actually operates; Section 5 reports a manipulation showing that this coverage is separable from optimization outcome, which is a caution against treating it as the mechanism.

Evaluation critiques and controlled decompositions. Each canonical audit in this genre pairs a critique with a reusable tool—a protocol, a stratified-bootstrap library, Nash averaging (Henderson et al. 2018; Agarwal et al. 2021; Balduzzi et al. 2018)—and within offline optimization SOO-Bench (Zhu, Qian et al. 2025) pursues a complementary benchmark-validity axis, stress-testing optimizer *stability* rather than attribution. The two-axis precedents are one-way fANOVA over model class at a single fixed search run (Hutter, Hoos, and Leyton-Brown 2014), a descriptive online-BO comparison that varies both axes without decomposing (Liang et al. 2021), and an HPO study that varies both and declines the decomposition by choice (Moosbauer et al. 2022). Our study is in this tradition but targets the surrogate $\times$ optimizer confound in offline MBO, which no member of it decomposes.

3 The Grid, and Five Confounds in It

The grid. Three surrogates $\times$ three acquisition optimizers $\times$ 7 synthetic tasks (Branin-2D through Griewank-30D, dataset drawn once at seed 0) $\times$ 30 seeds, plus 7 Design-Bench tasks at 16 seeds. **Surrogates:** a $K=5$ ensemble of two-hidden-layer ReLU MLPs (width 96, MSE, 35 epochs, Adam lr 3×10^{-3}), with μ, σ the member mean and standard deviation; an *exact GP* (BoTorch SingleTaskGP at its default covariance module, fit by marginal likelihood on a score-biased subsample, $N_{\text{max}}=800$); and a *sparse variational GP* (SVGP; 128 inducing points, ARD Matérn-5/2, 250 ELBO steps, $N_{\text{max}}=2000$). **Optimizers:** 100 Adam steps on x (lr 0.05, box-clipped); 5 rounds of Gaussian hill-climbing; and CMA-ES. All three surrogates are differentiable, so all nine cells are realizable. Every cell maximizes the LCB $\mu(x) - \beta\sigma(x)$ at $\beta=2$ and is scored by the ground-truth oracle on its 128 returned designs at the 100th percentile. Per task we min-max normalize the nine cell means and take a two-way ANOVA; η_{surr}^2 and η_{opt}^2 are the fractions of normalized-score variance carried by the two main

effects, with task-and-seed hierarchical bootstrap intervals ($B=10,000$).

Two details of that description are themselves engine state, which is the paper’s point in miniature: the exact GP’s default covariance is ARD RBF under the BoTorch pinned for the synthetic grid but ARD Matérn-5/2 under the older version stamped on the Design-Bench runs, and the GPs alone spend per-run marginal-likelihood budget, which a matched-tuning arm freezes to zero for every surrogate (the surrogate effect retains 76% under that control). Neither is recoverable from a method description that does not stamp its engine.

Against that grid we name five confounds. Each is a specific line of the scoring path, not a judgment call, and each carries a signed effect on the headline (Table 1).

Confound 1: target scaling. The ensemble regresses on raw targets spanning -2613 to $+36$ while both GP surrogates z -score theirs (`mbo.py:36–37, 130–138` against `mbo.py:255, 311–312`). Correcting this alone moves η_{surr}^2 from 0.367 [$0.254, 0.559$] to 0.283 [$0.186, 0.444$]. This is a measurement confound, not a bug report: the prior analysis states a min–max normalization that exists only in the analysis path and never in training, and the gap between stated and actual protocol is the finding.

Confound 2: candidate/oracle protocol. Two of three optimizers propose 256 designs and report the oracle-selected best 128, while the third proposes 128 (`mbo.py:384–389, 188, 198–201, 292, 392–394`). The reported percentile is therefore two different estimands in one column. Correcting this alone moves η_{surr}^2 from 0.367 to 0.450 [$0.312, 0.649$]. This is the strongest of the five: the protocol the prior analysis declares it holds identical is the identifiability license for the whole comparison, and the license is false.

Confound 3: ensemble size K . K enters as a confound to be controlled, never as a mechanism, and it has two halves that must be read together. It is *sensitive*: η_{surr}^2 is $0.326, 0.366, 0.408, 0.389$ at $K = 2, 3, 5, 10$, peaking at exactly the $K=5$ the field fixes by convention (Lakshminarayanan, Pritzel, and Blundell 2017). And it does *not* reverse: at $K=2$ the ensemble’s normalized cell marginal is 0.283 against the exact GP’s 0.767 and the SVGP’s 0.739 , so our own pre-registered kill criterion—“if the ensemble’s marginal at $K=2$ exceeds the GP’s, the surrogate main effect is a K artifact and the headline reverses”—did not fire. We also report that the figure this criterion was written against, a task-normalized ensemble score of 0.95 at $K=2$ in the prior supplement, *does not reproduce* on the audited engine, which gives 0.283 at that cell. (These two 0.283 s are unrelated quantities: Confound 1’s is a variance-explained statistic, this one is a normalized cell marginal. Neither corroborates the other.) Our sweep runs over $K \in \{2, 3, 5, 10\}$ and is therefore disjoint below the $K \in \{5, 10\}$ range over which ensemble surrogates were found robust (Li, Rudner, and Wilson 2024), and the decline with K runs against the direction reported for ensemble quality (Abe et al. 2022)—both further reasons to treat K as a nuisance coordinate rather than a result.

Table 1: The five confounds. Signed effect is on η_{surr}^2 (synthetic, $n=7$, 30 seeds) except Confound 5, which acts on η_{opt}^2 . The two *protocol* confounds move the headline against the ensemble-size confound; the net is Section 4.

Confound	What it breaks	Signed effect
1. Target scaling	ensemble trains on raw y	$0.367 \rightarrow 0.283$
2. Candidate/oracle	two estimands in one column	$0.367 \rightarrow 0.450$
3. Ensemble size K	headline fixed at the peak	$0.326-0.408$
4. β – σ match	per-task, not aggregate	ratio $0.07-1.44$
5. Query budget	optimizer = search effort	$0.005 \rightarrow 0.038$

Confound 4: β – σ mismatch. We registered that a shared $\beta=2$ delivers different effective pessimism to different surrogate classes, with the kill criterion “if the ratio is ~ 1 , the acquisition was matched and this confound does not exist.” *That kill fired.* The median $\sigma_{\text{GP}}/\sigma_{\text{ens}}$ is 1.19 at matched $K=5$ (1.21 pooled over K), so the fifth unmatched hyperparameter does not exist grid-wide, and we withdraw the aggregate claim. What survives is per-task: the ratio spans 0.07 on Branin-2D to 1.44 on Styblinski-5D and Rastrigin-15D, a 20-fold spread that a single shared multiplier cannot track. The general principle that a fixed multiplier on an uncalibrated interval is not comparable across model classes is owned by the interval-validity literature (Dewolf, De Baets, and Waegeman 2022) and its offline-RL instance (Ghasemipour, Gu, and Nachum 2022); we claim only the offline-MBO acquisition-comparison application, in its narrowed per-task form.

Confound 5: search-intensity budget. Surrogate-query budgets are unmatched by $11.8\times$ across the optimizer axis. The counts are *measured* by a counting proxy rather than derived: gradient $51,456$, perturbation $4,352$, CMA-ES median $6,528$ (range $932-6,536$). Three corrections belong here. The previously estimated “ $6\times-59\times$ ” spread was never measured, and CMA’s nominal budget of 3000 is a cap that rarely binds because the library stops on its own tolerances first. The proposal-count half of this confound is already removed by Confound 2, so the surviving inequality is in surrogate queries alone and does not double-count. And the trust region flagged in our own audit is inactive in the main grid, so equalizing queries isolates budget. Matching every optimizer at $Q=51,456$ (achieved $51,200 / 51,456 / 51,472$; 0 of 630 cells more than 5% off target) moves the optimizer main effect from a published η_{opt}^2 of 0.005 to 0.038 [$0.003, 0.123$]—an eightfold understatement, which we disclose. What that magnitude licenses is a separate question, deferred to Section 7.

4 The De-Confounding Protocol and Its Net Result

The protocol. Standardize the ensemble’s regression targets. Equalize the candidate proposal count across optimizers and impose one selection rule, so the oracle never chooses the reported set. Hold K , β and the query budget fixed and *report them*. Stamp the engine—library versions, flags, commit—




Figure 1: $\eta^2_{\text{surrogate}}$ against β (left, bootstrap 95% intervals at the measured endpoints) and against K (right), with the conventionally fixed operating point marked. The headline is reported at the maximum of both curves, and the ranking reverses nowhere on either.

with every number. Switching the first two on and off gives four corners (Table 2).

The net result runs against the genre’s direction. Correcting the two measurement and protocol confounds does not shrink the rewarded surrogate advantage but *strengthens* it, from a published η^2_{surr} of 0.37 to a corrected 0.405 [0.290, 0.556]; the ensemble-size confound moves it the other way, so the strengthening is the net of the two protocol confounds and not a property of de-confounding in general. The distinction matters because the broad claim is false. Audits of this genre usually shrink the effect they audit (Melis, Dyer, and Blunsom 2018), but not always: an ImageNet replication found a relative slope above one—1.69 on CIFAR-10, 1.11 on ImageNet, explicitly “the opposite of the standard overfitting scenario” (Recht et al. 2019)—and stratified reanalysis of deep-RL results moves effects both ways (Agarwal et al. 2021). A removed confound can as easily have been suppressing an effect as manufacturing one (Bressan 2019). Our claim is the narrow scalar one: we are not aware of an ML de-confounding audit that has produced a variance-explained statistic exceeding its own published value—unprecedented as an instance, not surprising as a statistic.

The decomposition is not resolvable at $n=7$ tasks, and we say so with every use of it. We pre-registered that it would not be, and the prediction confirmed. The four corner intervals have widths 0.26–0.34 against a corner range of only 0.167, and the corrected corner’s interval [0.290, 0.556] contains every other corner’s point estimate. So 0.405 never appears in this paper without its interval, and no corner is asserted to differ from another. The bootstrap is validated against the published precision: our uncorrected corner’s interval width, 0.305, reproduces the published 0.32.

The firmer result is a sensitivity, led by β . η^2_{surr} rises monotonically with pessimism—0.184 [0.085, 0.354], 0.214, 0.282, 0.406 [0.285, 0.564], 0.408 at $\beta = 0, 0.5, 1, 2, 5$ —and the $\beta=0$ and $\beta=2$ intervals barely overlap where the corner intervals do not separate at all (Figure 1). It is budget-dependent as well, more than doubling from 0.243 at $Q=4,352$ to 0.526 at $Q=51,456$. The honest reading is that the *magnitude* of any headline η^2_{surr} is a joint artifact of

Table 2: Four-corner decomposition (synthetic, 7 tasks, 30 seeds, $B=10,000$ bootstrap). X1 standardizes the ensemble’s targets; X3 equalizes the candidate/oracle protocol. The intervals overlap heavily: corner differences are *not* resolvable at $n=7$. Friedman is a 9-cell rank omnibus.

X1 / X3	η^2_{surr} [95% CI]	η^2_{opt}	η^2_{inter}	Friedman p
off / off	0.367 [0.254, 0.559]	0.013	0.165	6.1×10^{-5}
on / off	0.283 [0.186, 0.444]	0.036	0.146	1.7×10^{-3}
off / on	0.450 [0.312, 0.649]	0.006	0.152	4.1×10^{-5}
on / on	0.405 [0.290, 0.556]	0.005	0.160	8.1×10^{-4}

four chosen operating-point coordinates— K , β , the query budget, and the engine corner—so a headline that does not state all four is not reproducible. This strengthens the decomposition rather than weakening it, because the *direction* is invariant to all four: across every β , K , budget and corner we ran, the GP marginals sit at 0.75–0.85 and the ensemble’s at 0.24–0.36, never crossing. Only the size of the gap moves.

5 Scoped Mechanism: What the Advantage Is Not

Distance-aware uncertainty in GP-like models is established (Liu et al. 2020), as is the reading of a UCB-style acquisition as local search (Fan et al. 2024). What follows applies and diagnoses those results in the offline setting, and its main product is negative: five controls that could each have explained the GP–ensemble gap, and a sixth that could have confirmed the account we previously gave, all came back negative. For a paper whose contribution is a confound taxonomy, ruling out the tempting explanations is the work.

σ is a distance signal, not an error signal. The ensemble’s σ correlates with pointwise absolute error at only $\rho \approx 0.07$, but with k -NN distance to the training data at $\rho \approx 0.26$ —three to four times larger, positive on all seven tasks, and above 0.28 on the four mid- and high-dimensional ones. The prior analysis concluded σ was uninformative by measuring it against the wrong target; this is the corrected measurement, bounded by prior work on distance-aware uncertainty (Liu et al. 2020; van Amersfoort et al. 2020).

Elimination 1: it is not σ . Removing pessimism entirely ($\beta=0$, pure posterior-mean maximization) and reading both conditions on one β -invariant normalizer—a per-task min–max fit once over the pooled cell means of both conditions, refit inside each bootstrap resample—the GP–ensemble marginal gap is 0.319 [0.196, 0.460] at $\beta=0$ against 0.525 [0.406, 0.614] at $\beta=2$. Sixty-one percent of the advantage survives with σ gone, and the interval excludes zero. But the pessimism increment is itself significant—0.203 [0.007, 0.396], $p=0.020$ —so the correct statement is base-plus-amplification: pessimism *amplifies* a mean-quality base rather than creating it, and any claim that the advantage is independent of pessimism is refuted by our own data. A normalization-free cross-check (per-task z -score) gives 0.904 against 1.473 standard deviations—the same 0.61 ratio—so this is not an artifact of min–max scaling.

Our pre-registered collapse criterion passes less comfortably than earlier per- β figures implied: 0.319 against a threshold of 0.263, a margin inside the bootstrap interval.

Elimination 2: it is not finite width. The gap could be an under-parameterization artifact, since wide networks approach GP behaviour (Jacot, Gabriel, and Hongler 2018; Lee et al. 2018, 2019) and finite networks show spectral bias (Rahaman et al. 2019). Our K -sweep does *not* test that— K is cardinality, not per-member capacity, and offering it here would be a category error—so we ran the ablation that does. At fixed $K=5$, sweeping per-member width over a $10.7\times$ range leaves the gap unchanged: 0.480 [0.365, 0.576] at $w=96$ against 0.476 [0.208, 0.647] at $w=1024$, 99.1% retained, shrinkage -0.006 $[-0.210, 0.161]$. The curve is flat with noise and non-monotone—it dips to 0.336 at $w=256$ and returns—so it is not a decay. Two limits are ours to state. The interval widens monotonically with w ($0.211 \rightarrow 0.439$), making $w=1024$ the least precise point on the curve, so the supported claim is that the gap *does not close*, not that it is identical there. And the sweep stops at 1024 while the kernel limits are asymptotic, so this answers the objection at practical widths and makes no asymptotic claim. As a validity check, $w=96$ reproduces the incumbent grid bit-exactly for the gradient and CMA cells on all seven tasks.

Elimination 3: it is not predictive accuracy. The ensemble is the *more* accurate surrogate. Its held-out normalized RMSE on a 20% split never seen by the grid is lower than the GP’s at every width—0.445, 0.405, 0.393, 0.388 at $w = 96, 256, 512, 1024$ against the GP’s 0.479—winning in 26 of 28 (task, width) cells and on all seven tasks at $w \geq 512$ (six of seven at $w=96$, where Styblinski-5D is the exception). And it still loses the optimization comparison by roughly 0.48. The registered form of this test was narrower, and we report it as primary: read only at the two cells where ensemble and GP RMSE are statistically indistinguishable (Styblinski-5D at $w=256$, difference CI $[-0.021, 0.028]$, and at $w=512$, $[-0.043, 0.010]$), the mean gap is still 0.375. The 7/7 result is a post-hoc strengthening, labelled as one. Held-out NLL is not usable as a second accuracy axis—the GP’s is two orders of magnitude worse—because that is a calibration finding, orthogonal to this one. *The more accurate surrogate is the one that loses*, which removes predictive error from the candidate explanations.

Elimination 4: it is not budget, and more budget makes it worse. Equalizing surrogate queries does not close the gap; it does something more informative. η_{surr}^2 more than doubles with search budget, 0.243 at $Q=4,352$ to 0.526 at $Q=51,456$, and the composition of that rise is load-bearing: the ensemble’s marginal *falls* as budget rises ($0.361 \rightarrow 0.240$) while both GPs rise ($0.755 \rightarrow 0.849$ and $0.713 \rightarrow 0.794$). Give every optimizer more queries and the ensemble gets relatively worse—more search pressure finds more of the off-distribution maxima its mean admits. This was not what the budget arm was designed to test, so it is an observation rather than a pre-registered result, and corroborates a mechanism rather than any η^2 magnitude.

Eliminations 5 and 6: it is not how rough the mean is, and it is not premise coverage. Our previous account named the roughness of the surrogate’s mean as the axis. We tested that directly, bidirectionally, under pre-registration, and it failed. Penalizing the ensemble’s mean gradient cuts its measured roughness by up to 98.0% (spectral normalization, 90.3%)—a manipulation that landed by two orders of magnitude in the intended direction. The gap did not close; it *widened* for every variant (0.306 at base against 0.322, 0.346, 0.342 and 0.552), no closing interval excludes zero, and the Lipschitz-bounded ensemble is markedly worse in score ($0.657 \rightarrow 0.537$ at $w=96$, $\rightarrow 0.346$ at $w=1024$, where the constraint binds hardest). The analyzer is not rigged toward this outcome: a self-test planting synthetic ground truths confirms the same code returns the opposite verdict when smoothing genuinely closes a gap. Two observations sharpen the kill. Smoothing fixes the *premise* and not the outcome—own-proposal coverage rises monotonically to 0.998 under the strongest penalty with no effect on the gap—so the coverage proxy is separable from what it was supposed to explain, corroborating on a manipulated axis the weak $\rho(\text{coverage, score}) = 0.19$ we measure across cells and the dissociation in which the SVGP has the worst coverage of any GP (0.575) and the best score (0.881). And smoothing makes gradient collapse *worse*: the inversion rate rises from 0.548 to 0.695, a flatter mean giving ascent less to climb. The complementary arm, roughening the GP, is reportable only as void: no kernel setting raised the fitted mean’s roughness by the registered 25% (best +12.6%), because a posterior mean conditioned on ~ 800 points is smooth from the conditioning rather than the kernel, so roughness and fit quality are not independently manipulable in a fitted GP.

What survives. After six eliminations the surviving statement is narrow, and narrower than the one we previously published: what differs between the surrogate classes is *which* off-distribution maximizers their means admit, not how rough those means are. That is a diagnosis, not a mechanism with a causal test behind it, and we label it as such. A confound taxonomy earns its keep by removing the explanations a reader would otherwise reach for first, and these six are exactly those.

Where the collapse is genuine, and where it is not. Our own released tuning sweep concluded that the ensemble’s gradient collapse is an artifact a trust region repairs, and that result was absent from the prior analysis. We report it in full, including the corner where it fires against us. By the sweep’s own 5% rule, the count of tasks whose collapse survives tuning is 2/7 with neither correction applied, 0/7 under target scaling alone, 5/7 under the protocol correction alone, and 4/7 on the audited engine. The collapse is therefore genuine surrogate geometry on most tasks under the audited engine, driven by the candidate/oracle correction rather than target scaling—and on one corner it is not genuine at all.

Two framings we retract. First, the collapse is not gradient-specific. We registered that ensemble coverage would also be low on CMA proposals; it is. Under the matched protocol, gradient and CMA proposals differ in




Figure 2: Two eliminations in one image. At fixed $K=5$ the gap is flat with noise across a $10.7\times$ width range (intervals widen with w : the claim is that the gap does not close, not that it is identical), while the ensemble’s held-out normRMSE falls monotonically and overtakes the GP’s.

out-of-distribution coverage by 0.010, against 0.117 under the unmatched one, because both optimizers pool every iterate and return the top 128 by surrogate LCB. The claim narrows to the ensemble paired with any aggressive selector, and we withdraw the gradient-specific framing by name. Second, the premise-coverage value of 0.97 in the prior analysis belongs to a scikit-learn exact GP, not the differentiable GP the grid scores, whose coverage averages 0.831 and falls to 0.38 on Styblinski-5D; both are reported with the model each belongs to named. Coverage of $\{\mu - f \leq \beta\sigma\}$ is exactly the validity of the lower bound, the two events coinciding, which is what makes these frequencies worth measuring—and why the manipulation above matters, since a premise can be driven to 1.0 while the gap does not move.

6 Design-Bench: A Null, and Why It Is Null

The surrogate effect is at the floor, and the omnibus does not reject. Across all four engine corners the Design-Bench surrogate main effect is essentially zero— η_{surr}^2 of 0.001, 0.032, 0.002 and 0.018 on five tasks, every point estimate at the floor of its own bootstrap interval—and the Friedman omnibus fails to reject at $p = 0.19$ to 0.76 (Table 3). We therefore report no detectable difference at this power. We do not report equivalence, and the word does not appear in this section, because we are below the threshold for the test we ran: this omnibus is recommended for more than ten datasets (Demšar 2006), at $n=7$ a paired test needs $|d_z| \geq 1.27$ for 80% power at $\alpha=0.05$ —larger than anything we observe—and a failure to reject is not a demonstration that an effect is absent (Agarwal et al. 2021). Two robustness notes belong with the number. Dropping GFP leaves η_{surr}^2 at 0.001–0.021 with the omnibus still not rejecting; adding the two MuJoCo tasks raises it only to 0.044–0.091 but does make the omnibus reject in three of four corners ($p = 0.011$ to 0.016), and what it detects there is the *optimizer* axis, the subject of the mechanism below. We also disclose a pooling defect in the prior analysis: its normalization ranged over eleven cells, two of them arms the grid section never defines. Because mean-rank conclusions depend on the pool included (Benavoli, Corani, and Mangili 2016), we unify to the nine defined cells before any rank claim.

The null is not manufactured by approximate oracles. We pre-registered this competing mechanism and the test that would kill it. The null survives on the exact-oracle subset in every corner ($p = 0.34$ to 0.51), while the only rejections anywhere are on the random-forest-oracle subset ($p = 0.02$ to 0.05 in three of four corners). Repeated oracle evaluation gives a variance at machine precision ($\leq 2.3 \times 10^{-13}$, and 3.6×10^{-15} on four of five tasks), so the null is not noise drowning a signal. The counter-objection is ours to state: with two exact-oracle tasks this subset test has very little power, so the result is consistent with the null rather than a powered acceptance, corroborated by $\eta_{\text{surr}}^2 \approx 0$ and the noise floor rather than resting alone.

Some cells are frozen, and that explains the awkward number. Part of this null is structural rather than statistical. On TF-Bind-8, four of nine cells return exactly 1.0 on all 16 seeds with zero variance—the three exact-GP cells and ensemble-with-perturbation. That value is not a score: scores are min–max normalized to the offline dataset’s own range, whose maximum is attained by seven tied rows already in the data, so a cell reporting 1.0 returned a design no better than the best one it was given. On Ant, two of three GP cells return the bit-identical constant 1.5287419557571411 with standard deviation exactly 0.0 across all sixteen seeds, and a third (SVGP with perturbation) hits it on ten of sixteen. Three cells across two surrogate classes converging on one bit-identical value is retrieval, not coincidence.

The mechanism label follows, and it is the opposite of the one we had. Decode snap-back—movement in a relaxed logit space that argmax-decodes back to the starting design—predicts that an exact-constant freeze requires a decode step to snap back through. Ant is continuous with no argmax decode anywhere in its path, so our pre-registered kill against that account fires directly, and we report it as ours. The surviving mechanism is LCB paralysis: the GP’s lower confidence bound is locally maximal at the data, so the optimizer never leaves—a freeze occurring even without a decode step. This is the offline instance of a known reading of UCB-style acquisitions as local search (Fan et al. 2024), applied rather than discovered.

This buys an explanation for the section’s most awkward number. The optimizer axis *inverts* on Design-Bench—perturbation leads in all four corners (0.76–0.81 against gradient’s 0.40–0.68) and η_{opt}^2 runs two to three times η_{surr}^2 —while the synthetic grid calls the optimizer negligible. If the GP is frozen, every optimizer scores it identically, the surviving between-optimizer variance comes almost entirely from the unfrozen ensemble cells, and the axis inverts as a consequence: the Design-Bench inversion and the synthetic optimizer null become one finding seen through a frozen surrogate. Four limits are ours. It is two of three Ant GP cells, not three—the gradient cell returns 1.3242 ± 0.1975 and is not frozen. Fact and label are separable, so losing the label does not lose the frozen cells. Frozen cells and genuinely near-tied cells are reported separately, never pooled. And the optimizer half of this section stays provisional: budget matching ran on the synthetic tasks only, so Design-Bench carries Confound 5 unremoved on exactly the axis carrying

Table 3: Design-Bench: the surrogate main effect at the floor in every corner and task set (ranges over the four corners). Friedman p over the 9 defined cells. The omnibus rejects only on the 7-task set, where what it detects is the optimizer axis.

Task set	η_{sur}^2	η_{opt}^2	Friedman p
5 tasks	0.001–0.032	0.050–0.173	0.19–0.76
7 tasks (+MuJoCo)	0.044–0.091	0.096–0.197	0.011–0.157
4 tasks (−GFP)	0.001–0.021	0.047–0.184	0.44–0.67

the effect here, and we do not promote it on a synthetic-only arm.

7 Refuted Predictions, and What We Do Not Claim

Four of our own predictions came back against us. We report all four, because undisclosed they would read as selective reporting to anyone who opens the repository. (i) Our registered headline was that the acquisition optimizer explains most of the reported gap. It does not: the optimizer main effect is 0.038 even after correcting the budget imbalance suppressing it, and the pre-committed contingency—pivot to the mechanism rather than quietly restate—is what this paper does. (ii) Our ensemble-size criterion, “if the ensemble’s marginal at $K=2$ exceeds the GP’s, the surrogate main effect is a K artifact and the headline reverses,” did not fire, but its accompanying prediction confirmed: the headline is reported at the K that maximizes it. (iii) Our β – σ criterion, “if the ratio is ~ 1 , the acquisition was matched and this confound does not exist,” fired outright, and Confound 4 survives only in its narrowed per-task form. (iv) We registered that ensemble coverage would also be low on CMA proposals, retracting our gradient-specific framing. It was, and it did. The smoothness manipulation of Section 5 is a fifth and the most expensive: designed to confirm the mechanism we had published, it refuted it instead.

The optimizer magnitude, and what it licenses. Under a matched surrogate-query budget of $Q=51,456$, $\eta_{\text{opt}}^2 = 0.038$ [0.003, 0.123], so optimizer choice explains little of the variance in this grid. Three qualifications travel with it. It replaces a published 0.005 that understates by roughly eightfold, a correction we disclose rather than one put to us. The confirmation is not comfortable: the upper bound of 0.123 lies inside our pre-registered inconclusive band, so an effect up to about 0.12 is not excluded at $n=7$ —the same precision limit that makes the corner decomposition unresolvable. And the secondary budget level does not corroborate cleanly (0.066 [0.014, 0.340] at $Q=4,352$), so the null is established at high budget and underpowered at low, never unqualified. Most importantly, the *best* optimizer changes with budget: gradient leads at low budget (0.731) and perturbation at high (0.732), where the unmatched grid showed gradient modestly best throughout. An $11.8\times$ imbalance hides itself precisely by making one optimizer look mildly best everywhere. A small η_{opt}^2 therefore licenses “optimizer choice explains little vari-

ance” and never “optimizer choice is arbitrary”—the second reading is false on our own data.

What we do not claim. We make no field-reversal claim. That the choice of statistical model often matters more than the acquisition heuristic is the organizing doctrine of the standard Bayesian-optimization review (Shahriari et al. 2016) and is a decade old; it governs online BO and the choice of acquisition *function family*, not the numerical search routine in the offline setting. What our design falsifies is the local premise of one recent paper—that offline black-box optimization has focused on improving surrogate models while using fixed search strategies (Chemingui et al. 2024)—by holding one shared protocol and varying both axes. We make no claim about the field’s general beliefs about surrogates versus acquisitions. Nor do we claim ensemble size as a discovered mechanism (Li, Rudner, and Wilson 2024; Abe et al. 2022), equivalence on Design-Bench (Demšar 2006; Agarwal et al. 2021), or distance-aware uncertainty as ours (Liu et al. 2020; Fan et al. 2024).

Limitations. Seven synthetic tasks is a small n , and it binds the corner decomposition, the optimizer interval and the Design-Bench omnibus alike. The synthetic datasets are drawn once at seed 0, so reported variance is training and optimization variance, not data-draw variance. GFP is quarantined from coverage claims because its decode is degenerate, and we report that both ways: excluding it, one claim flips—the exact GP’s mean in-distribution coverage moves from below nominal to above it—while the effect-size null is unchanged. Budget matching covers the synthetic tasks only, which is why the optimizer half of Section 6 stays provisional. And the surviving mechanism statement is a diagnosis with six eliminations behind it and no positive causal test in front of it; the one causal test we ran refuted the account we had.

8 Conclusion

The reported advantage of GP-LCB in offline MBO is measured through five confounds specific to this comparison, and removing the two protocol confounds moves the surrogate main effect from a published 0.37 to 0.405 [0.290, 0.556]—the opposite direction from the one audits usually run, and the net of two corrections pushing against a third. That decomposition is not resolvable at seven tasks and we report it that way; the firmer result is that the magnitude of any $\eta_{\text{surrogate}}^2$ here is a joint artifact of four operating-point coordinates papers do not state. What the advantage *is* remains open, narrowed by elimination rather than demonstration: not σ , width, budget, held-out accuracy, mean roughness, or premise coverage—six explanations, each with a control behind it, the last refuting an account we had published. What we offer is the taxonomy, the protocol, and the engine stamp that makes either reproducible.

Ethical Statement

This work studies optimization methodology on public benchmarks and introduces no new data or deployed system.

References

- Abe, T.; Buchanan, E. K.; Pleiss, G.; Zemel, R.; and Cunningham, J. P. 2022. Deep Ensembles Work, But Are They Necessary? In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Agarwal, R.; Schwarzer, M.; Castro, P. S.; Courville, A.; and Bellemare, M. G. 2021. Deep Reinforcement Learning at the Edge of the Statistical Precipice. In *Advances in Neural Information Processing Systems*.
- Angelopoulos, A. N.; and Bates, S. 2023. Conformal Prediction: A Gentle Introduction. *Foundations and Trends in Machine Learning*, 16(4): 494–591.
- Balduzzi, D.; Tuyls, K.; Perolat, J.; and Graepel, T. 2018. Re-evaluating Evaluation. In *Advances in Neural Information Processing Systems*.
- Benavoli, A.; Corani, G.; and Mangili, F. 2016. Should We Really Use Post-Hoc Tests Based on Mean-Ranks? *Journal of Machine Learning Research*, 17(5): 1–10.
- Bressan, P. 2019. Confounds in “Failed” Replications. *Frontiers in Psychology*, 10: 1884.
- Brookes, D. H.; Park, H.; and Listgarten, J. 2019. Conditioning by Adaptive Sampling for Robust Design. In *International Conference on Machine Learning*, 773–782. PMLR.
- Chemingui, Y.; Deshwal, A.; Hoang, T. N.; and Doppa, J. R. 2024. Offline Model-Based Optimization via Policy-Guided Gradient Search. In *Proceedings of the AAAI Conference on Artificial Intelligence*.
- Demšar, J. 2006. Statistical Comparisons of Classifiers over Multiple Data Sets. *Journal of Machine Learning Research*, 7: 1–30.
- Dewolf, N.; De Baets, B.; and Waegeman, W. 2022. Valid Prediction Intervals for Regression Problems. *Artificial Intelligence Review*, 55: 577–613.
- Fan, Z.; Wang, W.; Ng, S. H.; and Hu, Q. 2024. Minimizing UCB: a Better Local Search Strategy in Local Bayesian Optimization. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Ferrari Dacrema, M.; Cremonesi, P.; and Jannach, D. 2019. Are we really making much progress? A worrying analysis of recent neural recommendation approaches. In *Proceedings of the 13th ACM Conference on Recommender Systems (RecSys)*, 101–109.
- Ghasemipour, K.; Gu, S. S.; and Nachum, O. 2022. Why So Pessimistic? Estimating Uncertainties for Offline RL through Ensembles, and Why Their Independence Matters. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Henderson, P.; Islam, R.; Bachman, P.; Pineau, J.; Precup, D.; and Meger, D. 2018. Deep Reinforcement Learning that Matters. In *Proceedings of the AAAI Conference on Artificial Intelligence*.
- Hutter, F.; Hoos, H.; and Leyton-Brown, K. 2014. An Efficient Approach for Assessing Hyperparameter Importance. In *International Conference on Machine Learning (ICML)*, 754–762.
- Jacot, A.; Gabriel, F.; and Hongler, C. 2018. Neural Tangent Kernel: Convergence and Generalization in Neural Networks. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Jin, Y.; Yang, Z.; and Wang, Z. 2021. Is Pessimism Provably Efficient for Offline Reinforcement Learning? In *International Conference on Machine Learning*, 5084–5096. PMLR.
- Kim, M.; Gu, J.; Yuan, Y.; Yun, T.; Liu, Z.; Bengio, Y.; and Chen, C. 2025. Offline Model-Based Optimization: Comprehensive Review. *arXiv preprint arXiv:2503.17286*.
- Krishnamoorthy, S.; Mashkaria, S. M.; and Grover, A. 2023. Diffusion Models for Black-Box Optimization. In *International Conference on Machine Learning*, 17842–17857. PMLR.
- Kumar, A.; Zhou, A.; Tucker, G.; and Levine, S. 2020. Conservative Q-Learning for Offline Reinforcement Learning. In *Advances in Neural Information Processing Systems*, volume 33, 1179–1191.
- Lakshminarayanan, B.; Pritzel, A.; and Blundell, C. 2017. Simple and Scalable Predictive Uncertainty Estimation using Deep Ensembles. In *Advances in Neural Information Processing Systems*, volume 30.
- Lee, J.; Bahri, Y.; Novak, R.; Schoenholz, S. S.; Pennington, J.; and Sohl-Dickstein, J. 2018. Deep Neural Networks as Gaussian Processes. In *International Conference on Learning Representations (ICLR)*.
- Lee, J.; Xiao, L.; Schoenholz, S. S.; Bahri, Y.; Novak, R.; Sohl-Dickstein, J.; and Pennington, J. 2019. Wide Neural Networks of Any Depth Evolve as Linear Models Under Gradient Descent. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Li, Y. L.; Rudner, T. G. J.; and Wilson, A. G. 2024. A Study of Bayesian Neural Network Surrogates for Bayesian Optimization. In *International Conference on Learning Representations*.
- Liang, Q.; Gongora, A. E.; Ren, Z.; Tiihonen, A.; et al. 2021. Benchmarking the Performance of Bayesian Optimization across Multiple Experimental Materials Science Domains. *npj Computational Materials*, 7: 188.
- Liu, J. Z.; Lin, Z.; Padhy, S.; Tran, D.; Bedrax-Weiss, T.; and Lakshminarayanan, B. 2020. Simple and Principled Uncertainty Estimation with Deterministic Deep Learning via Distance Awareness. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Lucic, M.; Kurach, K.; Michalski, M.; Gelly, S.; and Bousquet, O. 2018. Are GANs Created Equal? A Large-Scale Study. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Melis, G.; Dyer, C.; and Blunsom, P. 2018. On the State of the Art of Evaluation in Neural Language Models. In *International Conference on Learning Representations (ICLR)*.
- Moosbauer, J.; Binder, M.; Schneider, L.; Pfisterer, F.; Becker, M.; Lang, M.; Kotthoff, L.; and Bischl, B. 2022. Automated Benchmark-Driven Design and Explanation of Hyperparameter Optimizers. *IEEE Transactions on Evolutionary Computation*, 26(6): 1336–1350.

- Musgrave, K.; Belongie, S.; and Lim, S.-N. 2020. A Metric Learning Reality Check. In *European Conference on Computer Vision (ECCV)*, 681–699.
- Rahaman, N.; Baratin, A.; Arpit, D.; Draxler, F.; Lin, M.; Hamprecht, F.; Bengio, Y.; and Courville, A. 2019. On the Spectral Bias of Neural Networks. In *International Conference on Machine Learning (ICML)*, 5301–5310.
- Recht, B.; Roelofs, R.; Schmidt, L.; and Shankar, V. 2019. Do ImageNet Classifiers Generalize to ImageNet? In *International Conference on Machine Learning (ICML)*, 5389–5400.
- Shahriari, B.; Swersky, K.; Wang, Z.; Adams, R. P.; and de Freitas, N. 2016. Taking the Human Out of the Loop: A Review of Bayesian Optimization. *Proceedings of the IEEE*, 104(1): 148–175.
- Srinivas, N.; Krause, A.; Kakade, S. M.; and Seeger, M. W. 2012. Gaussian Process Optimization in the Bandit Setting: No Regret and Experimental Design. *IEEE Transactions on Information Theory*, 58(5): 3250–3265.
- Stanton, S.; Maddox, W.; and Wilson, A. G. 2023. Bayesian Optimization with Conformal Prediction Sets. In *International Conference on Artificial Intelligence and Statistics*. PMLR.
- Tan, R.-X.; Xue, K.; Lyu, S.-H.; Shang, H.; Wang, Y.; Wang, Y.; Fu, S.; and Qian, C. 2025. Offline Model-Based Optimization by Learning to Rank. In *International Conference on Learning Representations*.
- Trabucco, B.; Geng, X.; Kumar, A.; and Levine, S. 2022. Design-Bench: Benchmarks for Data-Driven Offline Model-Based Optimization. In *International Conference on Machine Learning*, 21658–21676. PMLR.
- Trabucco, B.; Kumar, A.; Geng, X.; and Levine, S. 2021. Conservative Objective Models for Effective Offline Model-Based Optimization. In *International Conference on Machine Learning*, 10358–10368. PMLR.
- van Amersfoort, J.; Smith, L.; Teh, Y. W.; and Gal, Y. 2020. Uncertainty Estimation Using a Single Deep Deterministic Neural Network. In *International Conference on Machine Learning (ICML)*, 9690–9700.
- Zhu, Y.; Qian, H.; et al. 2025. SOO-Bench: Benchmarks for Evaluating the Stability of Offline Black-Box Optimization. In *International Conference on Learning Representations*.